

Introduction to Parallel Computing

Frank Willmore
February 6, 2012

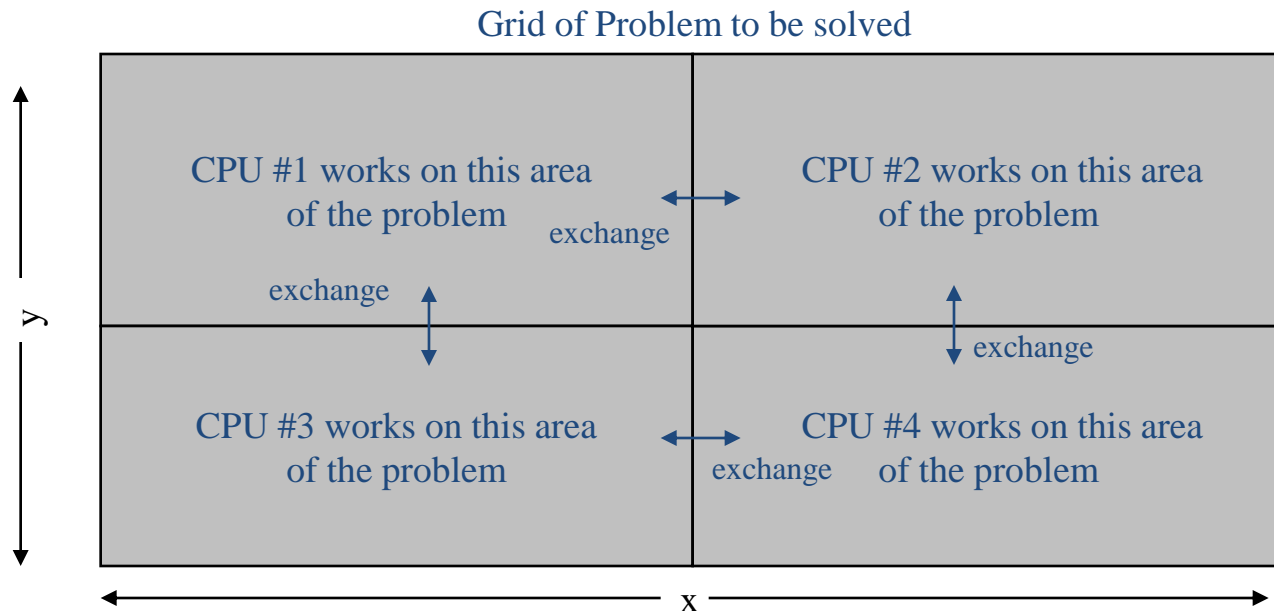
Outline

- Overview
- Theoretical background
- Parallel computing systems
- Parallel programming models
- MPI/OpenMP examples

OVERVIEW

What is Parallel Computing?

- Parallel computing: use of multiple processors or computers working together on a common task.
 - Each processor works on its section of the problem
 - Processors can exchange information



Why Do Parallel Computing?

- Limits of single CPU computing
 - performance
 - available memory
- Parallel computing allows one to:
 - solve problems that don't fit on a single CPU
 - solve problems that can't be solved in a reasonable time
- We can solve...
 - larger problems
 - the same problem faster
 - more cases
- All computers are parallel these days, even your iphone 4S has two cores...


THEORETICAL BACKGROUND

Speedup & Parallel Efficiency

- Speedup: $S_p = \frac{T_s}{T_p}$
 - p = # of processors
 - T_s = execution time of the sequential algorithm
 - T_p = execution time of the parallel algorithm with p processors
 - $S_p = p$ (linear speedup: ideal)

- Parallel efficiency

$$E_p = \frac{S_p}{p} = \frac{T_s}{pT_p}$$



Limits of Parallel Computing

- Theoretical Upper Limits
 - Amdahl's Law
 - Gustafson's Law
- Practical Limits
 - Load balancing
 - Non-computational sections
- Other Considerations
 - time to re-write code

Amdahl's Law

- All parallel programs contain:
 - parallel sections (we hope!)
 - serial sections (we despair!)
- Serial sections limit the parallel effectiveness
- Amdahl's Law states this formally
 - Effect of multiple processors on speed up

$$S_P \leq \frac{T_S}{T_P} \leq \frac{1}{f_s + \frac{f_p}{P}}$$


where

- f_s = serial fraction of code
- f_p = parallel fraction of code
- P = number of processors

Example:


$$f_s = 0.5, f_p = 0.5, P = 2$$
$$S_{p, max} = 1 / (0.5 + 0.25) = 1.333$$

Amdahl's Law



Practical Limits: Amdahl's Law vs. Reality

- In reality, the situation is even worse than predicted by Amdahl's Law due to:
 - Load balancing (waiting)
 - Scheduling (shared processors or memory)
 - Cost of Communications
 - I/O



Gustafson's Law


- Effect of multiple processors on run time of a problem with a *fixed amount of parallel work per processor*.

$$S_p \in P - \alpha \times (P - 1)$$

- α is the fraction of non-parallelized code where the parallel work per processor is fixed (not the same as f_p from Amdahl's)
- P is the number of processors

Comparison of Amdahl and Gustafson

Amdahl : fixed work




$$S \in \frac{1}{f_s + f_p / N}$$

$$S_2 \in \frac{1}{0.5 + 0.5 / 2} = 1.33$$

$$S_4 \in \frac{1}{0.5 + 0.5 / 4} = 1.6$$

Gustafson : fixed work per processor



$$S_p \in P - a \times (P - 1)$$

$$S_2 \in 2 - 0.5(2 - 1) = 1.5$$

$$S_4 \in 4 + 0.5(4 - 1) = 2.5$$

Scaling: Strong vs. Weak

- We want to know how quickly we can complete analysis on a particular data set by increasing the PE count
 - Amdahl's Law
 - Known as “strong scaling”
- We want to know if we can analyze more data in approximately the same amount of time by increasing the PE count
 - Gustafson's Law
 - Known as “weak scaling”

PARALLEL SYSTEMS

“Old school” hardware classification

	Single Instruction	Multiple Instruction
Single Data	SISD	MISD
Multiple Data	SIMD	MIMD

SISD No parallelism in either instruction or data streams (mainframes)

SIMD Exploit data parallelism (stream processors, GPUs)

MISD Multiple instructions operating on the same data stream. Unusual, mostly for fault-tolerance purposes (space shuttle flight computer)

MIMD Multiple instructions operating independently on multiple data streams (most modern general purpose computers, head nodes)

NOTE: GPU references frequently refer to SIMT, or single instruction multiple *thread*

Hardware in parallel computing


Memory access

- Shared memory
 - SGI Altix
 - IBM Power series nodes
- Distributed memory
 - Uniprocessor clusters
- Hybrid/Multi-processor clusters (Ranger, Lonestar)
- Flash based (e.g. Gordon)


Processor type

- Single core CPU
 - Intel Xeon (Prestonia, Wallatin)
 - AMD Opteron (Sledgehammer, Venus)
 - IBM POWER (3, 4)
- Multi-core CPU (since 2005)
 - Intel Xeon (Paxville, Woodcrest, Harpertown, Westmere, Sandy Bridge...)
 - AMD Opteron (Barcelona, Shanghai, Istanbul,...)
 - IBM POWER (5, 6...)
 - Fujitsu SPARC64 VIIIfx (8 cores)
- Accelerators
 - GPGPU
 - MIC

Shared and distributed memory




- All processors have access to a pool of shared memory
- Access times vary from CPU to CPU in NUMA systems
- Example: SGI Altix, IBM P5 nodes




- Memory is local to each processor
- Data exchange by message passing over a network
- Example: Clusters with single-socket blades

Hybrid systems




- A limited number, N , of processors have access to a common pool of shared memory
- To use more than N processors requires data exchange over a network
- Example: Cluster with multi-socket blades

Multi-core systems




- Extension of hybrid model
- Communication details increasingly complex
 - Cache access
 - Main memory access
 - Quick Path / Hyper Transport socket connections
 - Node to node connection via network

Accelerated (GPGPU and MIC) Systems



- Calculations made in both CPU and accelerator
- Provide abundance of low-cost flops
- Typically communicate over PCI-e bus
- Load balancing critical for performance

Accelerated (GPGPU and MIC) Systems



GPGPU (general purpose graphical processing unit)

- Derived from graphics hardware
- Requires a new programming model and specific libraries and compilers (CUDA, OpenCL)
- Newer GPUs support IEEE 754-2008 floating point standard
- Does not support flow control (handled by host thread)


MIC (Many Integrated Core)

- Derived from traditional CPU hardware
- Based on x86 instruction set
- Supports multiple programming models (OpenMP, MPI, OpenCL)
- Flow control can be handled on accelerator

Rendering a frame: Canonical example of a GPU task

- Single instruction: “Given a model and set of scene parameters...”
- Multiple data: Evenly spaced pixel locations (x_i, y_i)
- Output: “What are my red/green/blue/alpha values at (x_i, y_i) ?”
- The first uses of GPUs as accelerators were performed by posing physics problems as if they were rendering problems!

A GPGPU example:



Calculation of a free volume index over an evenly spaced set of points in a simulated sample of polydimethylsiloxane (PDMS)

- Relates directly to chemical potential via Widom insertion formalism of statistical mechanics
- Defined for all space
- Readily computable on GPU because of parallel nature of domain decomposition
- Generates voxel data which lends itself to spatial/shape analysis

$$FVI = e^{-\beta \psi_{ins}}$$

$$B_i = \frac{1}{V} \int FVI_i dV$$

$$\mu_i = -k_B T \ln \left(\frac{B_i}{\rho_i \lambda^3} \right)$$

PROGRAMMING MODELS

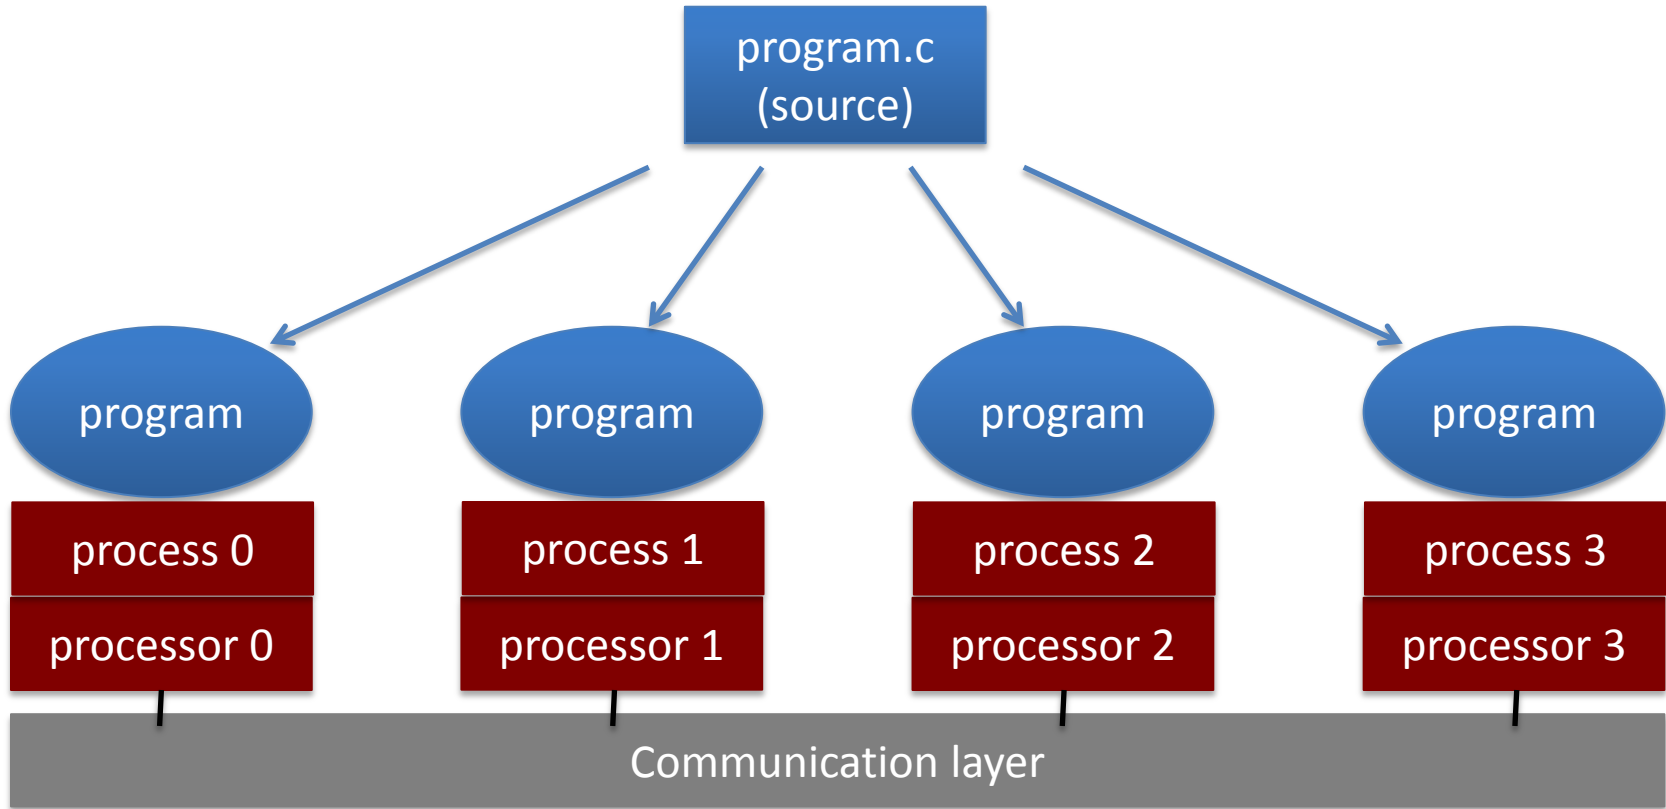
Types of parallelism

- Data Parallelism
 - Each processor performs the same task on different data (remember SIMD, MIMD)
- Task Parallelism
 - Each processor performs a different task on the same data (remember MISD, MIMD)
- Many applications incorporate both

Implementation: Single Program Multiple Data

- Dominant programming model for shared and distributed memory machines
- One source code is written
- Code can have conditional execution based on which processor is executing the copy
- All copies of code start simultaneously and communicate and synchronize with each other periodically

SPMD Model



Data Parallel Programming Example

- One code will run on 2 CPUs
- Program has array of data to be operated on by 2 CPUs so array is split into two parts.

CPU A

CPU B

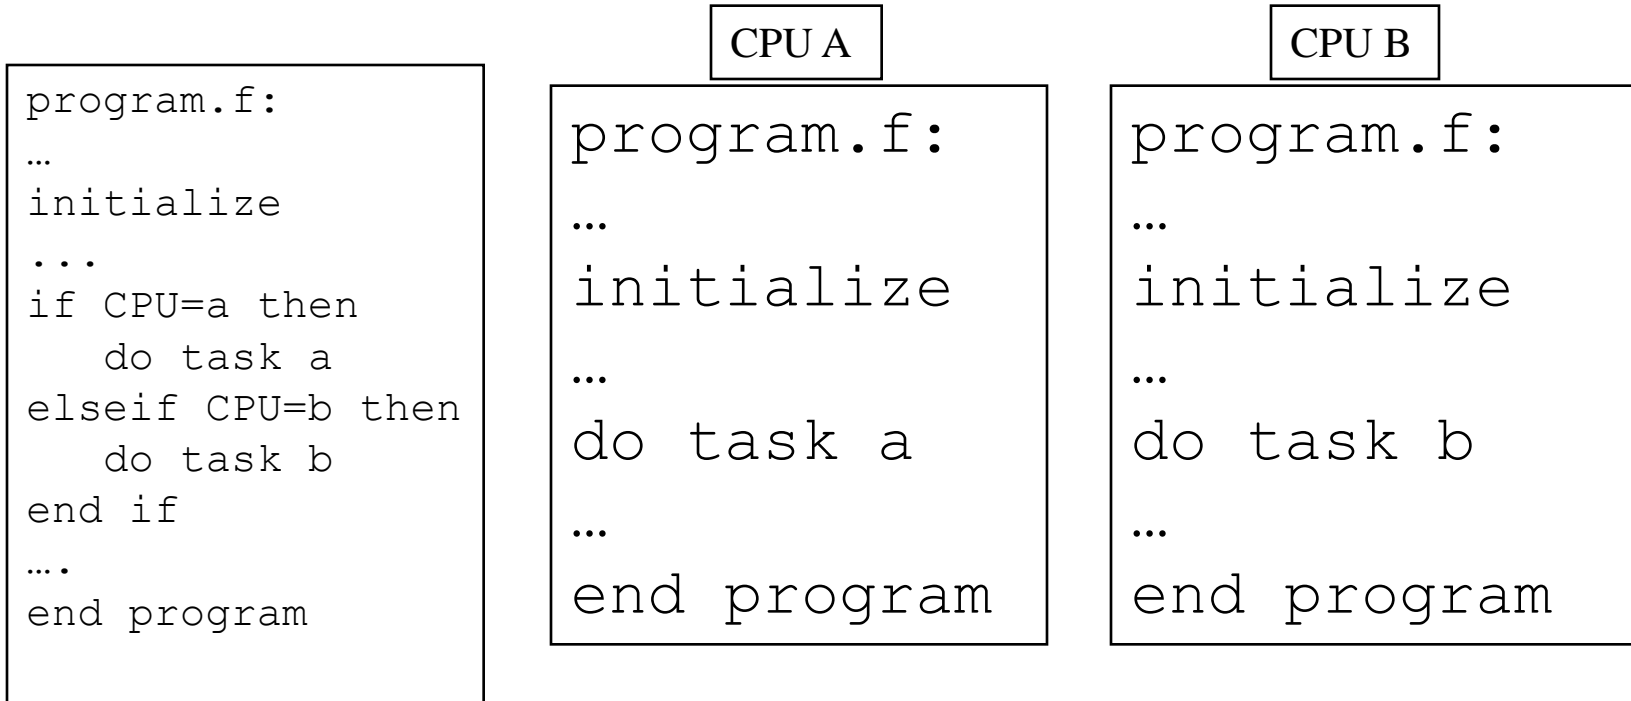
```
program:
...
if CPU=a then
    low_limit=1
    upper_limit=50
elseif CPU=b then
    low_limit=51
    upper_limit=100
end if
do I = low_limit,
upper_limit
    work on A(I)
end do
...
end program
```

```
program:
...
low_limit=1
upper_limit=50
do I= low_limit,
upper_limit
    work on A(I)
end do
...
end program
```

```
program:
...
low_limit=51
upper_limit=100
do I= low_limit,
upper_limit
    work on A(I)
end do
...
end program
```

Task Parallel Programming Example

- One code will run on 2 CPUs
- Program has 2 tasks (a and b) to be done by 2 CPUs



Shared Memory Programming: pthreads

- Shared memory systems (SMPs, ccNUMAs) have a single address space
- applications can be developed in which loop iterations (with no dependencies) are executed by different processors
- Threads are ‘lightweight processes’ (same PID)
- Allows ‘MIMD’ codes to execute in shared address space


Shared Memory Programming: OpenMP

- Built on top of pthreads
- shared memory codes are mostly data parallel, 'SIMD' kinds of codes
- OpenMP is a standard for shared memory programming (compiler directives)
- Vendors offer native compiler directives

Accessing Shared Variables

- If multiple processors want to write to a shared variable at the same time, there could be conflicts :

- Process 1 and 2
- read X
- compute $X+1$
- write X



- Programmer, language, and/or architecture must provide ways of resolving conflicts (mutexes and semaphores)

OpenMP Example #1: Parallel Loop

```
!$OMP PARALLEL DO  
  do i=1,128  
    b(i) = a(i) + c(i)  
  end do  
!$OMP END PARALLEL DO
```

- The first directive specifies that the loop immediately following should be executed in parallel.
- The second directive specifies the end of the parallel section (optional).
- For codes that spend the majority of their time executing the content of simple loops, the PARALLEL DO directive can result in significant parallel performance.

OpenMP Example #2: Private Variables

```
!$OMP PARALLEL DO SHARED(A,B,C,N) PRIVATE(I,TEMP)
do I=1,N
  TEMP = A(I)/B(I)
  C(I) = TEMP + SQRT(TEMP)
end do
!$OMP END PARALLEL DO
```

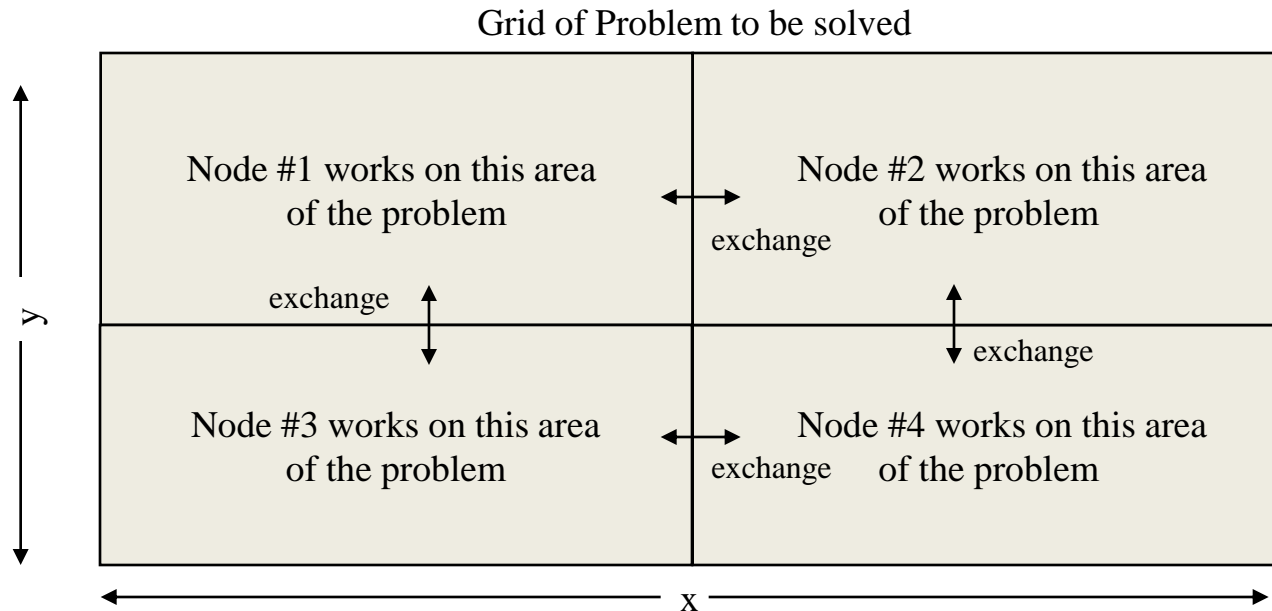
- In this loop, each processor needs its own private copy of the variable TEMP.
- If TEMP were shared, the result would be unpredictable since multiple processors would be writing to the same memory location.

Distributed Memory Programming: MPI

- Distributed memory systems have separate address spaces for each processor
- Local memory accessed faster than remote memory
- Data must be manually decomposed
- MPI is the de facto standard for distributed memory programming (library of subprogram calls)
- Vendors typically have native libraries such as SHMEM (T3E) and LAPI (IBM)

Data Decomposition

- For distributed memory systems, the 'whole' grid is decomposed to the individual nodes
 - Each node works on its section of the problem
 - Nodes can exchange information



Typical Data Decomposition

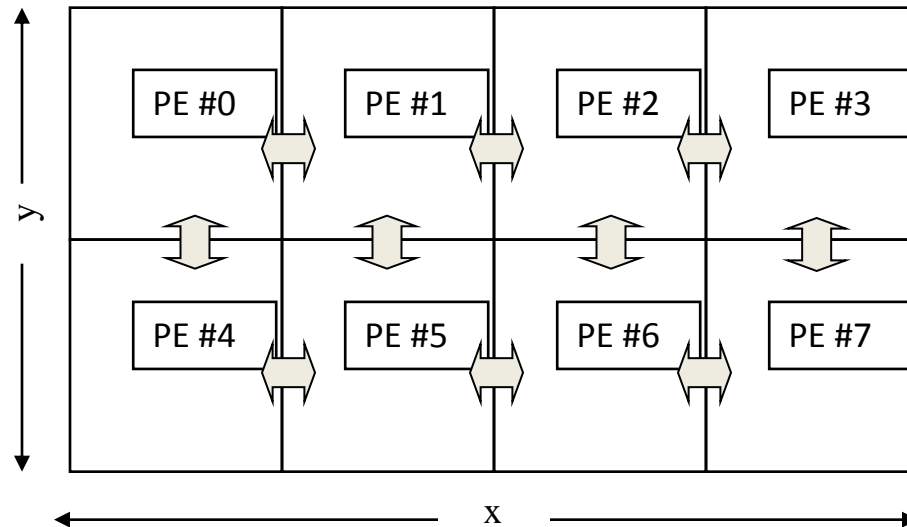
- Example: integrate 2-D propagation problem:

Starting partial
differential equation:

$$\frac{\partial \Psi}{\partial t} = D \cdot \frac{\partial^2 \Psi}{\partial x^2} + B \cdot \frac{\partial^2 \Psi}{\partial y^2}$$

Finite Difference
Approximation:

$$\frac{f_{i,j}^{n+1} - f_{i,j}^n}{\Delta t} = D \cdot \frac{f_{i+1,j}^n - 2f_{i,j}^n + f_{i-1,j}^n}{\Delta x^2} + B \cdot \frac{f_{i,j+1}^n - 2f_{i,j}^n + f_{i,j-1}^n}{\Delta y^2}$$



MPI Example #1

- Every MPI program needs these:

```
#include "mpi.h"
int main(int argc, char *argv[])
{
    int nPEs, iam;
    /* Initialize MPI */
    ierr = MPI_Init(&argc, &argv);
    /* How many total PEs are there */
    ierr = MPI_Comm_size(MPI_COMM_WORLD, &nPEs);
    /* What node am I (what is my rank?) */
    ierr = MPI_Comm_rank(MPI_COMM_WORLD, &iam);
    ...
    ierr = MPI_Finalize();
}
```

MPI Example #2

```
#include "mpi.h"


int main(int argc, char *argv[])
{
    int numprocs, myid;
    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &numprocs);
    MPI_Comm_rank(MPI_COMM_WORLD, &myid);
    /* print out my rank and this run's PE size */
    printf("Hello from %d of %d\n", myid, numprocs);
    MPI_Finalize();
}
```


MPI: Sends and Receives

- MPI programs must send and receive data between the processors (communication)
- The most basic calls in MPI (besides the three initialization and one finalization calls) are:
 - MPI_Send
 - MPI_Recv
- These calls are blocking: the source processor issuing the send/receive cannot move to the next statement until the target processor issues the matching receive/send.

Message Passing Communication

- Processes in message passing programs communicate by passing messages



- Basic message passing primitives: `MPI_CHAR`, `MPI_SHORT`, ...
- Send (parameters list)
- Receive (parameter list)
- Parameters depend on the library used
- Barriers

MPI Example #3: Send/Receive

```
#include "mpi.h"

int main(int argc, char *argv[])
{
    int numprocs, myid, tag, source, destination, count, buffer;
    MPI_Status status;

    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &numprocs);
    MPI_Comm_rank(MPI_COMM_WORLD, &myid);
    tag=1234;
    source=0;
    destination=1;
    count=1;

    if(myid == source){
        buffer=5678;
        MPI_Send(&buffer, count, MPI_INT, destination, tag, MPI_COMM_WORLD);
        printf("processor %d sent %d\n", myid, buffer);
    }
    if(myid == destination){
        MPI_Recv(&buffer, count, MPI_INT, source, tag, MPI_COMM_WORLD, &status);
        printf("processor %d got %d\n", myid, buffer);
    }
    MPI_Finalize();
}
```

Final Thoughts

- These are exciting and turbulent times in HPC.
- Systems with multiple shared memory nodes and multiple cores per node are the norm.
- Accelerators are rapidly gaining acceptance.
- Going forward, the most practical programming paradigms to learn are:
 - Pure MPI
 - MPI plus multithreading (OpenMP or pthreads)
 - Accelerator models (MPI or multithreading for MIC, CUDA or OpenCL for GPU)